

# RL for Adaptive Cloud Resource Allocation - Project Resources Guide

## Problem Definition (Why this is important)

Cloud providers (AWS, GCP, Azure, Kubernetes clusters, HPC centers) struggle with allocating resources (CPU, GPU, memory, network bandwidth) fairly and efficiently.

### Current Issues:

- Static allocation → over-provisioning (wastes money & energy) or under-provisioning (apps crash/slow)
- Reinforcement Learning (RL) can learn policies that dynamically adapt to workloads → auto-tuning

**Research Angle:** Most current solutions either focus only on throughput or only on fairness. You can propose a multi-objective RL scheduler (fairness + efficiency + energy-awareness).

## Feasibility Analysis

### 1. Data / Workloads ( Available)

You don't need access to AWS/GCP internals. Open workload traces exist:

- **Google Borg traces** (ClusterData 2011, 2019) → real jobs run on Google datacenters
- **Alibaba cluster traces** (2018) → CPU/memory usage patterns
- **HPC logs** from Parallel Workloads Archive (MIT, Hebrew University)

These traces can be replayed in a simulated Kubernetes/HPC cluster environment for testing your RL agent.

## 2. Simulation Environment ( Buildable)

### Options:

- Kubernetes + Minikube or SimGrid / CloudSim / OpenDC for cluster simulation
- Simpler approach: simulate workloads in Python → queue of jobs with resource demands, cluster with N nodes, RL agent allocates
- Eventually, you can deploy a prototype Kubernetes custom scheduler (written in Go or Python)

## 3. ML / RL Frameworks ( Manageable)

**Libraries:** Stable Baselines3, Ray RLlib, PyTorch RL

**Algorithms:** Start with DQN, PPO, A3C → extend to multi-agent RL (MARL) where each agent controls a cluster node

**Research Novelty:**

- Reward = weighted function of throughput + latency + fairness + energy cost
- Try hierarchical RL (global agent + node agents)

## 4. Division of Work (4 Members Team)

- **Person A:** Systems setup → Kubernetes/SimGrid + workload traces preprocessing
- **Person B:** RL model design + training pipeline (Stable Baselines3, PyTorch)
- **Person C:** Experimentation → benchmarking against baselines (First-Fit, Random, Heuristic schedulers)
- **Person D:** Results + Visualization + Research writing

This way, no one is idle and it scales well across a year.

## 5. Complexity Level

- **Engineering challenge:** Setting up simulation of workloads
- **Research challenge:** Designing a reward function that balances multiple objectives (throughput, fairness, energy)
- **Feasibility:** Definitely doable in 1 year, and you can push it towards publishable novelty

## Resources Needed

### 1. Hardware Resources

You don't need industrial-scale datacenters. A modest setup is enough for simulation-based research.

**Laptops/Desktops** (each team member's system can be used):

- **Minimum:** 8 GB RAM, i5/Ryzen processor
- **Recommended** (for training RL models): 1 system with GPU (NVIDIA GTX 1660 or better) → if not available, you can use Google Colab / Kaggle / cloud credits

**Optional: Cloud Credits** Many student programs give free credits:

- Google Cloud Student Credits (\$300 free)
- AWS Educate
- Microsoft Azure for Students (\$100 free)

You can deploy Kubernetes clusters on cloud VMs if needed.

**Feasibility:** Even if you don't get GPUs, you can simulate smaller workloads on CPU — RL training might be slower but still manageable.

## 2. Software / Tools

### (a) Cluster / Workload Simulation

**Option 1 (Realistic):**

- Kubernetes + Minikube (run on local system for cluster simulation)
- Workload submission with KubeFlow or custom scripts

### **Option 2 (Lightweight simulation):**

- CloudSim Plus (Java) or SimGrid (C++/Python)
- Or write a Python simulator: cluster = set of nodes, jobs = queue with CPU/memory demands

**Suggestion:** Start with Python simulator → once RL works, deploy it as a Kubernetes custom scheduler for real-world feel.

### **(b) Datasets (Workload Traces)**

Open-source workload traces to test your scheduler:

- **Google Cluster Data** (2011, 2019) → job CPU/memory demands
- **Alibaba Cluster Trace** (2018) → real production workloads
- **Parallel Workloads Archive (PWA)** → HPC traces (jobs, runtimes, nodes)

You can replay these traces in your simulator/Kubernetes environment.

### **(c) ML / RL Frameworks**

- Python (3.9+)
- PyTorch or TensorFlow (prefer PyTorch for flexibility)

- Stable Baselines3 (great for standard RL algorithms like PPO, DQN, A2C)
- Ray RLlib (scales well if you want to run multi-agent RL)

#### (d) Visualization & Metrics

- **Grafana + Prometheus** (if you use Kubernetes, for live monitoring)
- **Matplotlib, Seaborn** → plotting RL results

#### Metrics you'll need:

- Resource Utilization (CPU/mem)
- Job Completion Time / Makespan
- Fairness Index (Jain's index)
- Energy Consumption (can simulate via CPU utilization × power model)

### 3. Human Resources (Team Roles)

- **Team Member 1 (Systems Engineer):** Sets up Kubernetes/SimGrid, manages workload traces
- **Team Member 2 (RL Developer):** Implements RL agents in PyTorch/Stable Baselines3
- **Team Member 3 (Benchmarking & Analysis):** Runs experiments, compares with heuristics (First-Fit, Round-Robin, Default K8s scheduler)
- **Team Member 4 (Research & Writing):** Analyzes results, writes paper, handles visualization/metrics

## **4. Time Resources (Effort Breakdown)**

- **Learning curve** (2–3 months) → Kubernetes basics, RL basics
- **System Setup** (2 months) → simulator + baseline schedulers
- **RL Agent Implementation** (3–4 months)
- **Experiments & Tuning** (2 months)
- **Paper Writing & Finalization** (1–2 months)

## **5. Reference Material (to speed things up)**

### **Papers to start with:**

- "Deep Reinforcement Learning for Resource Scheduling in Cloud" (IEEE, 2020)
- "A Survey on Deep Reinforcement Learning for Resource Management in Cloud Computing" (ACM, 2021)

### **Books / Courses:**

- Sutton & Barto – Reinforcement Learning: An Introduction
- Kubernetes official docs (for scheduling)
- Stable Baselines3 tutorials

### **Expected Deliverables**

- **Simulator + RL Scheduler** (working code)
- **Comparison against baseline schedulers** (e.g., Kubernetes default scheduler)
- **Graphs & results:**
  - Resource utilization vs. job completion time
  - Fairness index (e.g., Jain's index)
  - Energy savings (optional)
- **Research paper** (~8–10 pages IEEE/ACM style)

## **Publication Targets (realistic for BTech level)**

- IEEE ICAC (International Conference on Autonomic Computing)
- ACM Middleware / IEEE CLOUD (student track)
- Springer LNCS / Elsevier FGCS (if aiming journal level)
- Or at least IEEE Xplore student conferences in India (ICCCNT, INDICON, TENCON, etc.)

## **Novelty Angles You Can Explore**

- Multi-Agent RL for fairness across nodes
- Carbon-aware scheduling → RL agent learns to allocate based on energy costs
- Adaptive reward shaping → RL balances short-term vs long-term gains
- Explainability in RL decisions (why a job got fewer resources)

## Final Verdict on Feasibility

With 4 motivated members + 1 year, this is very doable.

You'll have both:

- A working system demo (cluster + scheduler)
- Research novelty (multi-objective RL policy)

Strong enough for a paper publication + demo at project expo.

## Bottom Line on Resources

- **Hardware:** A few decent laptops + maybe cloud GPU credits
- **Software:** Kubernetes/SimGrid + Python (PyTorch + RL libraries)
- **Data:** Open workload traces (Google, Alibaba, PWA)
- **People:** Divide roles smartly, each member contributes to both coding + paper